

# Import Machine Learning Algorithms

```
In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from scipy import stats
import numpy as np
import plotly.express as px
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, roc_auc_score
from sklearn.metrics import roc_auc_score

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import RandomizedSearchCV
from xgboost import XGBClassifier as xgb

import shap
import lime
import lime.lime_tabular
```

```
In [2]: df = pd.read_csv("Iris Flower - Iris.csv")
df.head()
```

Out[2]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

```
In [3]: df = df.drop('Id',axis=1)
df.head()
```

Out[3]:

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

```
In [4]: df1 =df.copy()
df2 =df.copy()
```

```
In [5]: df.shape
```

Out[5]: (150, 5)

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   SepalLengthCm   150 non-null   float64
1   SepalWidthCm    150 non-null   float64
2   PetalLengthCm   150 non-null   float64
3   PetalWidthCm    150 non-null   float64
4   Species         150 non-null   object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
In [7]: df.describe()
```

```
Out[7]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

## The target Variable is PetalLengthCm

```
In [8]: df.rename(columns={'Species': 'Species_Value'}, inplace=True)  
df.head()
```

```
Out[8]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species_Value
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

## Chack Missing Values

```
In [9]: df.isnull().sum()
```

```
Out[9]: SepalLengthCm    0  
        SepalWidthCm     0  
        PetalLengthCm    0  
        PetalWidthCm     0  
        Species_Value    0  
        dtype: int64
```

```
In [10]: df.dtypes
```

```
Out[10]: SepalLengthCm    float64  
        SepalWidthCm     float64  
        PetalLengthCm    float64  
        PetalWidthCm     float64  
        Species_Value    object  
        dtype: object
```

## Data inconsistencies & Data Preprocessing:

```
In [11]: for col in df.columns:
          print(f"---##### {col} ---#####")
          print(df[col].value_counts())
```

Name: PetalLengthCm, dtype: int64

---##### PetalWidthCm ---#####

0.2	28
1.3	13
1.8	12
1.5	12
1.4	8
2.3	8
1.0	7
0.4	7
0.3	7
0.1	6
2.1	6
2.0	6
1.2	5
1.9	5
1.6	4
2.5	3
2.2	3
2.4	3

Handling NaN Values

```
In [12]: Num_cols = []
          cat_cols = []
          for col in df.columns:
              if df[col].dtypes == 'object':
                  cat_cols.append(col)
              else:
                  Num_cols.append(col)
```

```
In [13]: #handle null values of numerical columns
```

```
for col in Num_cols:  
    if df[col].isna:  
        df[col].fillna(df[col].median(), inplace=True)
```

```
In [14]: #handle null values of categorical columns
```


```
for col in cat_cols:  
    if df[col].isnull:  
        df[col].fillna(df[col].mode()[0], inplace=True)
```

```
In [15]: df.isnull().sum()
```

```
Out[15]: SepalLengthCm    0  
SepalWidthCm            0  
PetalLengthCm           0  
PetalWidthCm            0  
Species_Value           0  
dtype: int64
```

## Outliers

```
In [16]: plt.figure(figsize=(20,20))
for ax, col in enumerate(Num_cols):
    plt.subplot(5,6, int(ax+1))
    plt.title(col)
    sns.boxenplot(x=df[col],hue=df['Species_Value'])
plt.tight_layout()
plt.show()
```



**Encoding target value**

```
In [17]: from sklearn.preprocessing import LabelEncoder
led =LabelEncoder()
led.fit_transform(df['Species_Value'])
df['Species_Value'] =led.fit_transform(df['Species_Value'])
df.head()
```

Out[17]:

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species_Value
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

```
In [18]: df.value_counts('Species_Value')
```

```
Out[18]: Species_Value
0      50
1      50
2      50
dtype: int64
```

## Model\_Selection of ML



```
In [19]: df.head()
```

```
Out[19]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species_Value
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

## Decision Tree Classifier

```
In [20]: from sklearn import tree
```

```
In [21]: df.head()
```

```
Out[21]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species_Value
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

```
In [22]: x =df.drop('Species_Value',axis=1)
y =df['Species_Value']
```

## DecisionTreeClassifier

```
In [23]: from sklearn.tree import DecisionTreeClassifier
```

```
In [24]: dtc = DecisionTreeClassifier()
```

```
In [25]: dtc.fit(x,y)
```

```
Out[25]: ▾ DecisionTreeClassifier  
DecisionTreeClassifier()
```

```
In [26]: dtc.predict([[1,0,0,0]])
```

X does not have valid feature names, but DecisionTreeClassifier was fitted with feature names

```
Out[26]: array([0])
```

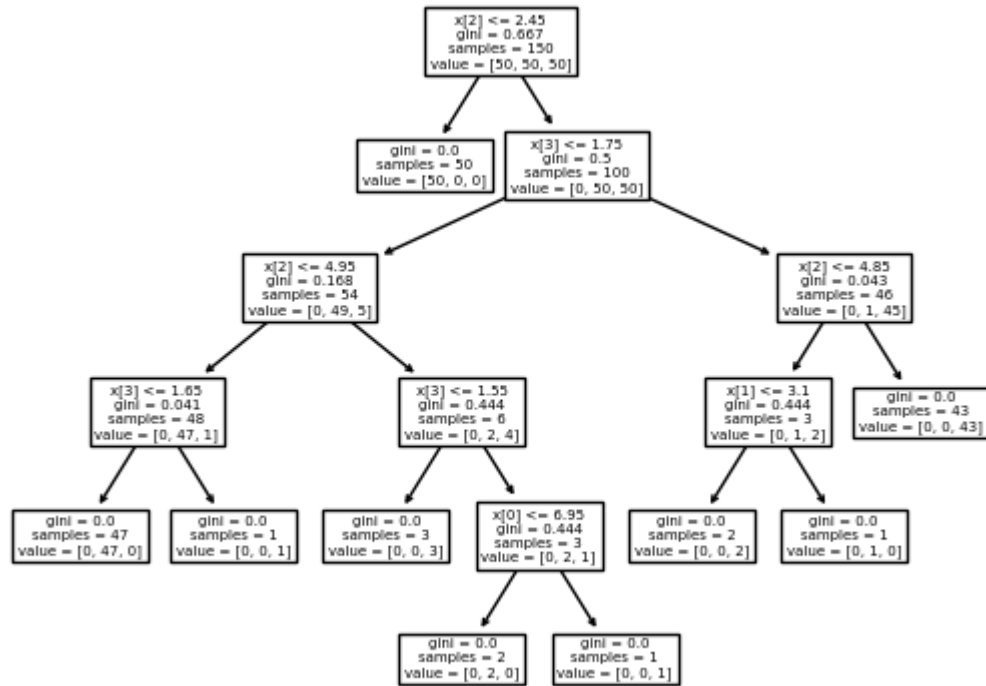
```
In [27]: dtc.predict([[0,0,0,1]])
```

X does not have valid feature names, but DecisionTreeClassifier was fitted with feature names

```
Out[27]: array([0])
```

```
In [28]: tree.plot_tree(dtc)
```

```
Out[28]: [Text(0.5, 0.9166666666666666, 'x[2] <= 2.45\ngini = 0.667\nsamples = 150\nvalue = [50, 50, 50]'),
Text(0.4230769230769231, 0.75, 'gini = 0.0\nsamples = 50\nvalue = [50, 0, 0]'),
Text(0.5769230769230769, 0.75, 'x[3] <= 1.75\ngini = 0.5\nsamples = 100\nvalue = [0, 50, 50]'),
Text(0.3076923076923077, 0.5833333333333334, 'x[2] <= 4.95\ngini = 0.168\nsamples = 54\nvalue = [0, 4
9, 5]'),
Text(0.15384615384615385, 0.4166666666666667, 'x[3] <= 1.65\ngini = 0.041\nsamples = 48\nvalue = [0, 4
7, 1]'),
Text(0.07692307692307693, 0.25, 'gini = 0.0\nsamples = 47\nvalue = [0, 47, 0]'),
Text(0.23076923076923078, 0.25, 'gini = 0.0\nsamples = 1\nvalue = [0, 0, 1]'),
Text(0.46153846153846156, 0.4166666666666667, 'x[3] <= 1.55\ngini = 0.444\nsamples = 6\nvalue = [0, 2,
4]'),
Text(0.38461538461538464, 0.25, 'gini = 0.0\nsamples = 3\nvalue = [0, 0, 3]'),
Text(0.5384615384615384, 0.25, 'x[0] <= 6.95\ngini = 0.444\nsamples = 3\nvalue = [0, 2, 1]'),
Text(0.46153846153846156, 0.08333333333333333, 'gini = 0.0\nsamples = 2\nvalue = [0, 2, 0]'),
Text(0.6153846153846154, 0.08333333333333333, 'gini = 0.0\nsamples = 1\nvalue = [0, 0, 1]'),
Text(0.8461538461538461, 0.5833333333333334, 'x[2] <= 4.85\ngini = 0.043\nsamples = 46\nvalue = [0, 1,
45]'),
Text(0.7692307692307693, 0.4166666666666667, 'x[1] <= 3.1\ngini = 0.444\nsamples = 3\nvalue = [0, 1,
2]'),
Text(0.6923076923076923, 0.25, 'gini = 0.0\nsamples = 2\nvalue = [0, 0, 2]'),
Text(0.8461538461538461, 0.25, 'gini = 0.0\nsamples = 1\nvalue = [0, 1, 0]'),
Text(0.9230769230769231, 0.4166666666666667, 'gini = 0.0\nsamples = 43\nvalue = [0, 0, 43]')]
```



```
In [29]: tree.plot_tree(dtc, rounded=True, filled=True, feature_names = x.columns )
```

```
Out[29]: [Text(0.5, 0.9166666666666666, 'PetalLengthCm <= 2.45\ngini = 0.667\nsamples = 150\nvalue = [50, 50, 50]'),
Text(0.4230769230769231, 0.75, 'gini = 0.0\nsamples = 50\nvalue = [50, 0, 0]'),
Text(0.5769230769230769, 0.75, 'PetalWidthCm <= 1.75\ngini = 0.5\nsamples = 100\nvalue = [0, 50, 50]'),
Text(0.3076923076923077, 0.5833333333333334, 'PetalLengthCm <= 4.95\ngini = 0.168\nsamples = 54\nvalue = [0, 49, 5]'),
Text(0.15384615384615385, 0.4166666666666667, 'PetalWidthCm <= 1.65\ngini = 0.041\nsamples = 48\nvalue = [0, 47, 1]'),
Text(0.07692307692307693, 0.25, 'gini = 0.0\nsamples = 47\nvalue = [0, 47, 0]'),
Text(0.23076923076923078, 0.25, 'gini = 0.0\nsamples = 1\nvalue = [0, 0, 1]'),
Text(0.46153846153846156, 0.4166666666666667, 'PetalWidthCm <= 1.55\ngini = 0.444\nsamples = 6\nvalue = [0, 2, 4]'),
Text(0.38461538461538464, 0.25, 'gini = 0.0\nsamples = 3\nvalue = [0, 0, 3]'),
Text(0.5384615384615384, 0.25, 'SepalLengthCm <= 6.95\ngini = 0.444\nsamples = 3\nvalue = [0, 2, 1]'),
Text(0.46153846153846156, 0.08333333333333333, 'gini = 0.0\nsamples = 2\nvalue = [0, 2, 0]'),
Text(0.6153846153846154, 0.08333333333333333, 'gini = 0.0\nsamples = 1\nvalue = [0, 0, 1]'),
Text(0.8461538461538461, 0.5833333333333334, 'PetalLengthCm <= 4.85\ngini = 0.043\nsamples = 46\nvalue = [0, 1, 45]'),
Text(0.7692307692307693, 0.4166666666666667, 'SepalWidthCm <= 3.1\ngini = 0.444\nsamples = 3\nvalue = [0, 1, 2]'),
Text(0.6923076923076923, 0.25, 'gini = 0.0\nsamples = 2\nvalue = [0, 0, 2]'),
Text(0.8461538461538461, 0.25, 'gini = 0.0\nsamples = 1\nvalue = [0, 1, 0]'),
Text(0.9230769230769231, 0.4166666666666667, 'gini = 0.0\nsamples = 43\nvalue = [0, 0, 43]')]
```



## K-Nearest Neighbors (KNN)

```
In [30]: from sklearn.neighbors import KNeighborsRegressor
```

```
In [31]: knn = KNeighborsRegressor(n_neighbors=3)
```

```
In [32]: knn.fit(df[["Petal.Length.Cm"]],df["Species_Value"])
```

```
Out[32]: KNeighborsRegressor
KNeighborsRegressor(n_neighbors=3)
```


```
In [33]: knn.predict([[15]])
```

X does not have valid feature names, but KNeighborsRegressor was fitted with feature names

```
Out[33]: array([2.])
```

```
In [34]: plt.scatter(df.PetalWidthCm, df.Species_Value)  
plt.scatter(15,2, color="red")
```

```
Out[34]: <matplotlib.collections.PathCollection at 0x20a4a5de2d0>
```



## Logistic Regression

```
In [35]: df2.head()
```

```
Out[35]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

```
In [36]: df2.value_counts('Species')
```

```
Out[36]: Species
Iris-setosa      50
Iris-versicolor  50
Iris-virginica   50
dtype: int64
```

```
In [37]: df2['Species'] = df2['Species'].replace(['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'],[0,1,2])
```

```
In [38]: df2.head()
```

```
Out[38]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0





```
In [47]: reg.predict_proba(x)
```

```
Out[47]: array([[9.81804276e-01, 1.81957100e-02, 1.43460484e-08],
 [9.71802160e-01, 2.81978107e-02, 2.97679380e-08],
 [9.85495978e-01, 1.45040094e-02, 1.21748649e-08],
 [9.76424964e-01, 2.35749969e-02, 3.91821962e-08],
 [9.85398356e-01, 1.46016317e-02, 1.18853668e-08],
 [9.70381261e-01, 2.96186651e-02, 7.36250444e-08],
 [9.86930605e-01, 1.30693757e-02, 1.97677029e-08],
 [9.76445341e-01, 2.35546320e-02, 2.74313737e-08],
 [9.79977769e-01, 2.00222013e-02, 3.01443225e-08],
 [9.69264678e-01, 3.07352902e-02, 3.12962705e-08],
 [9.76465697e-01, 2.35342841e-02, 1.92046470e-08],
 [9.75509897e-01, 2.44900592e-02, 4.34927557e-08],
 [9.74675031e-01, 2.53249478e-02, 2.11985210e-08],
 [9.92023518e-01, 7.97647851e-03, 3.82662375e-09],
 [9.88115480e-01, 1.18845174e-02, 2.82138058e-09],
 [9.86674116e-01, 1.33258709e-02, 1.29085668e-08],
 [9.88048548e-01, 1.19514428e-02, 9.20169859e-09],
 [9.81540546e-01, 1.84594348e-02, 1.95775987e-08],
 [9.56425841e-01, 4.35740903e-02, 6.85582969e-08],
 [9.84109627e-01, 1.58903526e-02, 2.04864764e-08],
 [9.46788727e-01, 5.32111874e-02, 8.60349466e-08],
 [9.81718341e-01, 1.82816258e-02, 3.27551617e-08],
 [9.96012118e-01, 3.98788090e-03, 1.30120720e-09],
 [9.52307406e-01, 4.76923584e-02, 2.35294370e-07],
 [9.52142739e-01, 4.78570566e-02, 2.04706858e-07],
 [9.51724920e-01, 4.82749940e-02, 8.58539787e-08],
 [9.69642321e-01, 3.03575937e-02, 8.57525714e-08],
 [9.74938912e-01, 2.50610632e-02, 2.48321606e-08],
 [9.77345878e-01, 2.26541046e-02, 1.73006443e-08],
 [9.71377478e-01, 2.86224638e-02, 5.78685246e-08],
 [9.64456751e-01, 3.55431792e-02, 6.96055318e-08],
 [9.64847917e-01, 3.51520261e-02, 5.72363014e-08],
 [9.88381497e-01, 1.16184958e-02, 7.03827774e-09],
 [9.89023113e-01, 1.09768821e-02, 5.32307431e-09],
 [9.69264678e-01, 3.07352902e-02, 3.12962705e-08],
 [9.84675735e-01, 1.53242569e-02, 7.90857558e-09],
 [9.78888779e-01, 2.11112114e-02, 9.59440027e-09],
 [9.69264678e-01, 3.07352902e-02, 3.12962705e-08],
 [9.85936993e-01, 1.40629917e-02, 1.52884991e-08],
 [9.74139741e-01, 2.58602304e-02, 2.82359399e-08],
 [9.86622239e-01, 1.33777494e-02, 1.12894387e-08],
 [9.62454256e-01, 3.75456797e-02, 6.46692339e-08],
 [9.89077657e-01, 1.09223321e-02, 1.11248863e-08],
```

[9.72415189e-01, 2.75846748e-02, 1.36473924e-07],  
[9.60295084e-01, 3.97046941e-02, 2.22313454e-07],  
[9.73941209e-01, 2.60587512e-02, 3.94699041e-08],  
[9.80339423e-01, 1.96605521e-02, 2.52583405e-08],  
[9.83423725e-01, 1.65762546e-02, 1.98937643e-08],  
[9.78568486e-01, 2.14314958e-02, 1.86534512e-08],  
[9.78710783e-01, 2.12891974e-02, 1.91086885e-08],  
[2.10862668e-03, 8.73911720e-01, 1.23979654e-01],  
[5.76823060e-03, 8.59784431e-01, 1.34447338e-01],  
[1.05081764e-03, 7.25004743e-01, 2.73944440e-01],  
[1.54491926e-02, 9.39605856e-01, 4.49449509e-02],  
[2.36330137e-03, 8.15138141e-01, 1.82498557e-01],  
[6.96159722e-03, 8.60065758e-01, 1.32972645e-01],  
[3.73034137e-03, 7.16634303e-01, 2.79635356e-01],  
[1.48642786e-01, 8.48302043e-01, 3.05517120e-03],  
[2.76433824e-03, 8.96565673e-01, 1.00669988e-01],  
[4.14097366e-02, 9.11964289e-01, 4.66259747e-02],  
[5.63081930e-02, 9.37213785e-01, 6.47802175e-03],  
[1.50880019e-02, 8.98923592e-01, 8.59884062e-02],  
[9.14621362e-03, 9.76491713e-01, 1.43620734e-02],  
[3.03185326e-03, 7.79195995e-01, 2.17772151e-01],  
[7.43534625e-02, 9.15191810e-01, 1.04547279e-02],  
[5.24771741e-03, 9.26352106e-01, 6.84001765e-02],  
[8.65632963e-03, 7.74751407e-01, 2.16592263e-01],  
[1.64699180e-02, 9.65138181e-01, 1.83919005e-02],  
[1.80985911e-03, 8.00836957e-01, 1.97353184e-01],  
[2.40284995e-02, 9.59352815e-01, 1.66186855e-02],  
[2.27943755e-03, 4.40397579e-01, 5.57322984e-01],  
[1.67935233e-02, 9.56795838e-01, 2.64106382e-02],  
[7.12108592e-04, 5.96065484e-01, 4.03222407e-01],  
[3.03114082e-03, 8.59773974e-01, 1.37194885e-01],  
[7.04928838e-03, 9.42941928e-01, 5.00087837e-02],  
[5.04702753e-03, 9.20090125e-01, 7.48628470e-02],  
[1.11365773e-03, 8.01376342e-01, 1.97510000e-01],  
[5.72573050e-04, 4.81147340e-01, 5.18280087e-01],  
[5.45016575e-03, 8.13094564e-01, 1.81455270e-01],  
[6.19098607e-02, 9.34706775e-01, 3.38336387e-03],  
[2.92562226e-02, 9.57114156e-01, 1.36296213e-02],  
[3.73398416e-02, 9.55117128e-01, 7.54303014e-03],  
[2.51785108e-02, 9.56466417e-01, 1.83550721e-02],  
[4.46077656e-04, 3.49641519e-01, 6.49912404e-01],  
[1.01618347e-02, 7.50991852e-01, 2.38846313e-01],  
[9.88064442e-03, 7.88903270e-01, 2.01216085e-01],

[2.24498796e-03, 8.05140196e-01, 1.92614816e-01],  
[2.76574182e-03, 9.13049949e-01, 8.41843090e-02],  
[2.69596896e-02, 9.28603489e-01, 4.44368212e-02],  
[1.99216288e-02, 9.38039653e-01, 4.20387180e-02],  
[8.71455579e-03, 8.97810703e-01, 9.34747407e-02],  
[4.61617435e-03, 8.28235221e-01, 1.67148604e-01],  
[1.75871757e-02, 9.56979931e-01, 2.54328934e-02],  
[1.22509266e-01, 8.74440260e-01, 3.05047361e-03],  
[1.44435519e-02, 9.20449668e-01, 6.51067802e-02],  
[1.99267378e-02, 9.38131157e-01, 4.19421055e-02],  
[1.70447215e-02, 9.25454598e-01, 5.75006806e-02],  
[8.46623899e-03, 9.35114317e-01, 5.64194440e-02],  
[2.44641374e-01, 7.54068127e-01, 1.29049945e-03],  
[1.91187811e-02, 9.36056233e-01, 4.48249863e-02],  
[8.83945976e-07, 3.92380755e-03, 9.96075309e-01],  
[2.40478186e-04, 1.62637592e-01, 8.37121929e-01],  
[2.43989458e-06, 2.55900308e-02, 9.74407529e-01],  
[3.08462316e-05, 8.17010657e-02, 9.18268088e-01],  
[3.67341536e-06, 1.74640166e-02, 9.82532310e-01],  
[5.42581633e-08, 4.64029239e-03, 9.95359653e-01],  
[5.74744241e-03, 5.13824824e-01, 4.80427733e-01],  
[6.12087217e-07, 2.13369896e-02, 9.78662398e-01],  
[5.17296248e-06, 5.32566160e-02, 9.46738211e-01],  
[6.38000895e-07, 5.75609365e-03, 9.94243268e-01],  
[2.97721865e-04, 2.10640717e-01, 7.89061562e-01],  
[7.19637222e-05, 1.37323809e-01, 8.62604227e-01],  
[2.09423186e-05, 6.53123413e-02, 9.34666716e-01],  
[2.27360977e-04, 1.45309102e-01, 8.54463537e-01],  
[6.80627890e-05, 4.35826418e-02, 9.56349295e-01],  
[5.07181450e-05, 5.41253766e-02, 9.45823905e-01],  
[5.48969877e-05, 1.22934156e-01, 8.77010947e-01],  
[8.24635299e-08, 3.56625143e-03, 9.96433666e-01],  
[3.09054643e-09, 1.00046829e-03, 9.98999529e-01],  
[3.88669705e-04, 4.51638362e-01, 5.47972968e-01],  
[5.49480546e-06, 2.38827666e-02, 9.76111739e-01],  
[6.07239782e-04, 1.90609152e-01, 8.08783608e-01],  
[3.08193167e-08, 4.65327551e-03, 9.95346694e-01],  
[5.80434319e-04, 3.93056939e-01, 6.06362627e-01],  
[1.26038590e-05, 3.86515075e-02, 9.61335889e-01],  
[4.77771806e-06, 5.14939827e-02, 9.48501240e-01],  
[1.06269412e-03, 4.56521122e-01, 5.42416184e-01],  
[1.01199819e-03, 3.85479491e-01, 6.13508511e-01],  
[1.05051565e-05, 3.63647722e-02, 9.63624723e-01],

```
[1.66741235e-05, 1.41905719e-01, 8.58077606e-01],
[1.04799974e-06, 2.91884542e-02, 9.70810498e-01],
[6.86347607e-07, 1.74357944e-02, 9.82563519e-01],
[7.76839814e-06, 2.72882629e-02, 9.72703969e-01],
[5.24875293e-04, 4.75421735e-01, 5.24053390e-01],
[6.22412534e-05, 1.88447131e-01, 8.11490627e-01],
[3.85676613e-07, 1.17528221e-02, 9.88246792e-01],
[1.13862839e-05, 1.73678116e-02, 9.82620802e-01],
[6.67619232e-05, 1.19535901e-01, 8.80397337e-01],
[1.60407457e-03, 4.40522497e-01, 5.57873429e-01],
[3.90465598e-05, 9.35554232e-02, 9.06405530e-01],
[6.18835283e-06, 2.03228444e-02, 9.79670967e-01],
[9.80090824e-05, 1.20781073e-01, 8.79120918e-01],
[2.40478186e-04, 1.62637592e-01, 8.37121929e-01],
[2.00921845e-06, 1.26057079e-02, 9.87392283e-01],
[3.73036310e-06, 1.21304884e-02, 9.87865781e-01],
[5.50264541e-05, 8.02057758e-02, 9.19739198e-01],
[2.25169948e-04, 2.51893935e-01, 7.47880895e-01],
[1.36503465e-04, 1.57222283e-01, 8.42641214e-01],
[4.47860511e-05, 3.85093811e-02, 9.61445833e-01],
[4.70348664e-04, 2.34999252e-01, 7.64530399e-01]])
```

```
In [48]: xtrain, xtest, ytrain, ytest = train_test_split(x,y, test_size=.3, random_state=42)
```

## LogisticRegression

```
In [49]: LR_model = LogisticRegression()
```

```
In [50]: LR_model.fit(xtrain, ytrain)
```

A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
Out[50]:
```

```
▼ LogisticRegression
LogisticRegression()
```

```
In [51]: smote = SMOTE(random_state=42)
xtrain, ytrain = smote.fit_resample(xtrain, ytrain)
```

## RandomForestClassifier

```
In [52]: RFC = RandomForestClassifier()
```

```
In [53]: RFC.fit(xtrain, ytrain)
```

A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

```
Out[53]: ▾ RandomForestClassifier
RandomForestClassifier()
```

## XGboost

```
In [54]: XGB = xgb()
```

```
In [55]: XGB.fit(xtrain, ytrain)
```

```
Out[55]: XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              multi_strategy=None, n_estimators=None, n_jobs=None,
```

## SVC

```
In [56]: svc = SVC(kernel='rbf', gamma='scale', probability=True)
```

```
In [57]: svc.fit(xtrain, ytrain)
```

A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
Out[57]: SVC
SVC(probability=True)
```

## Predicting and Evaluation Metrics



```
In [58]: trained_models = [  
    LR_model,  
    RFC,  
    XGB,  
    svc  
]
```

```
In [59]: for model in trained_models:
    y_pred = model.predict(xtest)
    accuracy = accuracy_score(ytest, y_pred)
    precision = precision_score(ytest, y_pred, average='micro')
    recall = recall_score(ytest, y_pred, average='micro')
    f1 = f1_score(ytest, y_pred, average='micro')
    cm = confusion_matrix(ytest, y_pred)

    proba = model.predict_proba(xtest)
    auc = roc_auc_score(ytest, proba, multi_class='ovr')


    print(f"\t ***** {model.__class__.__name__} *****")
    print("Accuracy Score: ", accuracy)
    print("Precision Score: ", precision)
    print("Recall Score: ", recall)
    print("F1 Score: ", f1)
    print("AUC Score: ", auc)
    print("Confusion Matrix: \n", cm)
    sns.heatmap(cm)
    plt.show()

    fpr, tpr, threshold = roc_curve(ytest, proba[:,1], pos_label=1)
    plt.plot(fpr, tpr, linestyle="--", label="CURVE", )
    plt.title("ROC CURVE")
    plt.xlabel("False Positive Rate")
    plt.ylabel("True Positive Rate")
    plt.legend()
    plt.show()
```



```
In [60]: sns.histplot(x = "PetalLengthCm", data = df2, kde=True, hue="Species")
```

```
Out[60]: <Axes: xlabel='PetalLengthCm', ylabel='Count'>
```



RandomForest Classifier has the highest accuracy. So we will use Gradient Boosting Classifier for hyper parameter tuning

Hyperparameter Tuning:

```
In [61]: params = {  
    'n_estimators': [100, 200, 300],  
    'criterion': ['gini', 'entropy'],  
    'max_depth': [None, 10, 20, 30],  
    'min_samples_split': [2, 5, 10],  
    'min_samples_leaf': [1, 2, 3],  
    'max_features': ['auto', 'sqrt', 'log2']  
}
```

```
In [62]: model_param = {
    'svm': {
        'model': SVC(gamma='auto'),
        'params': {
            'C': [1.0, 5.0, 10.0],
            'kernel': ['rbf', 'linear']
        }
    },
    'LogReg': {
        'model': LogisticRegression(solver='liblinear'),
        'params': {
            'C': [1.0, 5.0, 10.0],
            'penalty': ['l1', 'l2'],
        }
    },
    'rf': {
        'model': RandomForestClassifier(),
        'params': {
            'n_estimators': [20, 50, 100],
            'criterion': ['gini', 'entropy'],
            'min_samples_leaf': [1, 2],
            'max_features': ['sqrt', 'log2']
        }
    },
    'XGboost': {
        'model': xgb(),
        'params': {
            'learning_rate': [0.1, 0.01, 0.2],
            'n_estimators': [20, 50, 100],
        }
    }
}
```

```
In [64]: scores = []

for name, mp in model_param.items():
    RandomSearch = RandomizedSearchCV(estimator=mp['model'] , param_distributions=mp['params'], return_tr
    RandomSearch.fit(xtrain, ytrain)
    scores.append({
        'model': mp['model'],
        'best_score': RandomSearch.best_score_,
        'best_param': RandomSearch.best_params_
    })
```

[illegible]

```
In [65]: scores
```

```
Out[65]: [{'model': SVC(gamma='auto'),
  'best_score': 0.9636363636363636,
  'best_param': {'kernel': 'linear', 'C': 5.0}},
 {'model': LogisticRegression(solver='liblinear'),
  'best_score': 0.9549407114624506,
  'best_param': {'penalty': 'l1', 'C': 10.0}},
 {'model': RandomForestClassifier(),
  'best_score': 0.9458498023715414,
  'best_param': {'n_estimators': 100,
  'min_samples_leaf': 1,
  'max_features': 'sqrt',
  'criterion': 'entropy'}},
 {'model': XGBClassifier(base_score=None, booster=None, callbacks=None,
  colsample_bylevel=None, colsample_bynode=None,
  colsample_bytree=None, device=None, early_stopping_rounds=None,
  enable_categorical=False, eval_metric=None, feature_types=None,
  gamma=None, grow_policy=None, importance_type=None,
  interaction_constraints=None, learning_rate=None, max_bin=None,
  max_cat_threshold=None, max_cat_to_onehot=None,
  max_delta_step=None, max_depth=None, max_leaves=None,
  min_child_weight=None, missing=nan, monotone_constraints=None,
  multi_strategy=None, n_estimators=None, n_jobs=None,
  num_parallel_tree=None, random_state=None, ...),
  'best_score': 0.9276679841897234,
  'best_param': {'n_estimators': 20, 'learning_rate': 0.2}}]
```

Predict testing values with hyperparameter tuning

```
In [67]: rf = RandomForestClassifier(n_estimators=50, min_samples_leaf=1, criterion='entropy')
rf.fit(xtrain, ytrain)
b_predicted = rf.predict(xtest)
```

A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_sample s,), for example using ravel().



```
In [69]: accuracy_score(ytest, b_predicted)
```

```
Out[69]: 1.0
```


```
In [71]: accuracy = accuracy_score(ytest, b_predicted)
precision = precision_score(ytest, b_predicted, average='micro')
recall = recall_score(ytest, b_predicted, average='micro')
f1 = f1_score(ytest, b_predicted, average='micro')
cm = confusion_matrix(ytest, b_predicted)
```

```
In [72]: print("Accuracy Score: ", accuracy)
print("Precision Score: ", precision)
print("Recall Score: ", recall)
print("F1 Score: ", f1)
print("AUC Score: ", auc)
print("Confusion Matrix: \n", cm)
```

```
Accuracy Score: 1.0
Precision Score: 1.0
Recall Score: 1.0
F1 Score: 1.0
AUC Score: 1.0
Confusion Matrix:
[[19  0  0]
 [ 0 13  0]
 [ 0  0 13]]
```


```
In [73]: sns.heatmap(cm)
```

```
Out[73]: <Axes: >
```



Interpretability:

```
In [74]: explainer = shap.Explainer(rf)
shap_values = explainer.shap_values(xtest)
shap.summary_plot(shap_values, xtest)
```



```
In [ ]:
```